

Complex networks

Ts Evans

To cite this article: Ts Evans (2004) Complex networks, Contemporary Physics, 45:6, 455-474, DOI: [10.1080/00107510412331283531](https://doi.org/10.1080/00107510412331283531)

To link to this article: <https://doi.org/10.1080/00107510412331283531>



Published online: 08 Nov 2010.



Submit your article to this journal [↗](#)



Article views: 257



View related articles [↗](#)



Citing articles: 2 View citing articles [↗](#)

Complex networks

T. S. Evans

An outline of recent work on complex networks is given from the point of view of a physicist. Motivation, achievements and goals are discussed with some of the typical applications from a wide range of academic fields. An introduction to the relevant literature and useful resources is also given.

1. Introduction

These days, to gauge what is hot and what is not in the world of physics, one need only turn to the electronic preprint archives. If you look at the one focussed on condensed matter physics, cond-mat [1], you will notice that among the preprints looking at high temperature superconductivity, Bose–Einstein condensation and other traditional pursuits of this community, there are a large number talking about networks. Search for this keyword in the title and you will be completely overwhelmed by papers. What you may also notice is that these papers only started appearing in numbers from about 1998 as figure 1 shows. Scan some of these papers and you will quickly notice that these are not full of plots of conductivity and the like, but are counting links on web pages [2] or discussing data on the frequency of words in English texts [3, 4]. If you are naturally skeptical about trendy new areas of physics and attempts to mix physics with anything and everything, then the citations of papers in journals of sociology [5, 6] and of books on archeology and anthropology [7–9] may just be the last straw! However, one cannot deny that an awful lot of physicists have found something new and intriguing about networks over the last few years. What I hope to do in the article is to give a flavour of what the excitement is all about, and perhaps why underneath the hyperbolae accompanying any new developments, there are sensible questions waiting to be answered.

First, what are these networks? In the simplest form we are discussing a collection of points or *vertices* which are connected by a variety of lines or *edges*. In more complicated cases, we can add more information to our network. For instance if it represents a transport network there may be distances, times and/or capacities associated with our edges (*weighted edges*). Perhaps flow is only

possible in one direction along an edge so we think of them as having a direction (*directed edges*). Vertices may represent different types of location, perhaps factories, warehouses and shops (*coloured vertices*). However, while such extra data is important in many cases, networks are all fundamentally just a set of vertices and a set of edges connecting some of the vertex pairs. Given that such a simple concept is at the heart of many physical structures, it should not be surprising that mathematicians have been studying them for some time. What I have been calling a network is what is called a *graph* by mathematicians and I will use the terms network and graph interchangeably. Unfortunately, the terminology of graph theory is not standardized so one must be careful to check each author's definitions and to specify one's own notation (refer to table 1 for this work).

Once a physical network has been represented as a mathematical graph, then numerous problems can be considered. Optimization questions are often classic pro-

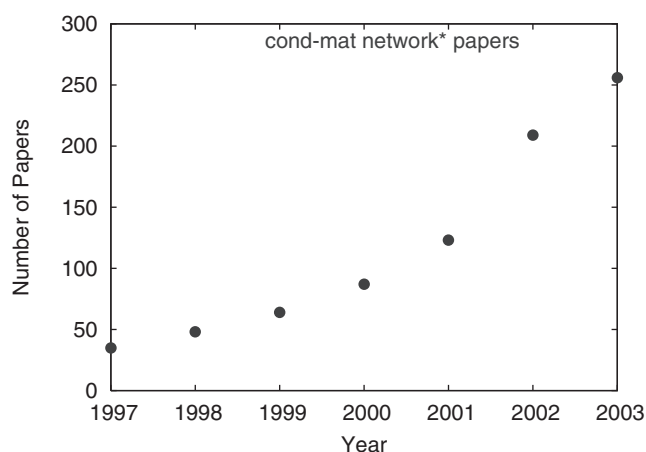


Figure 1. The number of papers listed per year on the cond-mat archives [1] with a word starting with 'network' in their title.

Author's address: Theoretical Physics, Blackett Laboratory, Imperial College London, Prince Consort Road, London, SW7 2BW, UK.

Table 1. Summary of various symbols used throughout the text. See text for definitions.

Symbol	Meaning
N	Number of vertices in a network
E	Number of edges in a network
k	Degree of a vertex
K	Average degree of vertices in a network
$n(k)$	Number of vertices of degree k in a network
L	Average of the shortest distances between all vertex pairs in a network
c	Clustering coefficient of a vertex

blems of graph theory. One example is the ‘travelling salesman’ problem in which one must find the shortest distance a salesman can travel given they are to visit a set of cities (the vertices) along prescribed routes between cities (the edges). Critical path analysis is another application of graph theory where one tries to identify bottlenecks in a process. However, the recent interest of physicists stems from the discovery of new ways of classifying and generating graphs and linking these new types to networks we see in the world around us. Applications include various types of human interaction such as social or business relationships. The internet is a very fashionable and often overworked topic, but it is a natural application for all that. In fact sociologists and anthropologists had been studying some of these systems for some time, sometimes using graph theory. However, physicists certainly provided new tools and views even if the debate predated them.

So let me now turn to look at the different types of network and how they have been used to study mathematical and physical problems.

2. Global network properties

I will focus on the most basic properties of a network, and hence I will ignore any directions or weights associated with edges, and any colours or labels added to vertices. Further, I will restrict myself to the case where multiple edges between the same pair of vertices are not allowed, and an edge is not allowed to start and end on the same vertex. These networks are sometimes called *simple graphs* although the terminology varies. It will also turn out that most networks of interest are *sparse*, that is the number of edges E is of the same order as the number of vertices N . This is considerably smaller than the potential $N(N-1)/2$ edges of a simple network.

As a physicist, the first type of network that comes to my mind is the regular lattice. Playing a fundamental role in solids, they are characterized by invariance under translation in space by a lattice spacing along a lattice axis¹. In this

case the vertices of a network could represent the atoms of a crystal. The edges could then indicate the most important interactions. For instance, in a simple two-dimensional regular square lattice, such as figure 2, each vertex is attached to four edges so that the *degree* of every vertex is four to use the graph theory terminology. It therefore has $E = 2N$ edges when there are N vertices, and so it is a sparse network.

In terms of abstract graph theory, it is easy to picture the opposite of a regular lattice. For instance if we took N vertices we could select at random E of the $N(N-1)/2$ different vertex pairs and connect them to give what is called a *random graph*². If E is of the same order of magnitude as N then we have a sparse network again. For instance the random graph of the same number of vertices and edges as the regular lattice of figure 2, given in figure 3, is sparse. The random graph immediately highlights one point central in the study of networks, namely that in general the vertices of a network have no position associated with them. Clearly in the case of the random network, the connections between vertices are made without any reference to notions of relative position of vertices unlike in the crystal example, and so we do not need to, and must not, picture these vertices as having any position. It is therefore wrong to think of the random graph as living in any particular dimension. Only the limitations of our visualization skills forces me to display the random graph of figure 3 in two dimensions, and the implicit coordinates of the vertices in the picture have no meaning. I could move them anywhere on the page and, provided I maintained the edges connecting the vertices, it would represent the same network. The central idea of graph theory and the issues studied with modern networks is that it is only the connections between vertices that are fundamental and these and their associated properties should form the starting point of any analysis.

Of course, the random graph of figure 3 has no structure other than its connections, as this was all that was used in its construction. The positions of vertices in some real space are not an issue. However, the physical properties of a material whose atoms are arranged in a square lattice do require us to take account of the physical coordinates so in such a case we must use the representation on the left of figure 2 and not those on the centre and right where the positional information has been discarded to a greater or lesser extent.

Problems with real crystals are not going to be advanced by studying networks in themselves. However, there are many problems where it is the connections themselves, and not their nature, which are the key. In this case the positions

¹Strictly speaking, only if they are infinite or periodic.

²Similarly, we could connect each edge with probability $p = 2E/(N(N-1))$ and obtain E edges on average. I will confine the term random network to these definitions although others sometimes refer to these as an Erdős–Reyni type random graph and extend the name to cover a wider family of graphs formed by some probabilistic rule.




Figure 2. A 20 node square periodic lattice, with 40 edges between nearest neighbours only, but shown in three different ways. In the left hand picture (a), edges going off the sides wrap around. In the second two examples, the spatial coordinates associated with each vertex are ignored. In the middle picture (b), the vertices are still displayed in a systematic order so the regular nature of the lattice is still visible. In the last picture (c) the vertices are displayed around the circle in a random order and the graph now looks random, even though in terms of simple graphs it is as regular as the other two. For later reference note that this lattice has an average distance of $L = 2.32$, diameter 4, and the average clustering coefficient is 0.




Figure 3. On the left a standard 20 node, 40 edge random graph. On the right a similar graph but constrained so that all vertices are attached to four edges, i.e. their degree is always 4 like the lattice earlier. For later reference, network A has average distance $L = 2.17$ (ignoring the disconnected vertex), diameter 5 and average clustering coefficient of $c = 0.134$. Network B has average shortest distance of $L = 2.22$, diameter 4 and average clustering coefficient of $c = 0.15$.

of the vertices in our world are not important. Vertices could be people, the links indicating a relationship which is not just a simple function of the physical distance between people. The ‘space’ of such anthropological networks is not the simple Euclidean three-dimensional one of most scientific problems. For instance, suppose we require a system of radio communication between sites (the vertices). An edge linking two sites represents a communication link and we will suppose that this requires a dedicated frequency that no other pair of sites can use. Perhaps for the system, the physical Euclidean distance between sites is not relevant, being small enough not to effect error rates etc. On the other hand, if the total bandwidth available is limited, then we only have a limited number of frequencies available, we can

only have a certain number of edges in our network. If the error rate in the communications rises with the number of edges traversed by each message, then we need to find a network that minimizes the average number of edges traversed in moving from one vertex to another on the graph, a purely graph-based concept.

Let us now look at the various ideas and concepts that can be defined with only the fundamental properties of a network. In other words, we can only work with a set of vertices and a set of edges joining some of these vertex pairs. We can imagine walking from vertex to vertex on this network, moving only along the edges of the network as in figure 4. In doing so we define paths on the networks between vertices. In some networks there may be no path at




Figure 4. Graph (a) on the left is connected since one can walk along the network from any vertex to any vertex. The distances from vertex v_3 to any other vertex are given in square brackets. So vertex v_6 is 3 from v_3 , the length of the shortest path via v_2 and v_5 (paths via v_4 for instance are longer and are not considered). This is also the diameter of this network as no pair of vertices has a shortest path longer than 3. The average of shortest distances over all 30 pairs of vertices is $L = 1.87$. Removing the v_2 – v_5 edge leaves the disconnected graph (b) on the right since one cannot then walk for instance from v_1 to v_6 . The v_1, v_2, v_3, v_4 vertices and their edges form one connected component of the right-hand graph.

all between some vertex pairs. This means that the graph is *disconnected* and appears in two or more distinct pieces called *components*.

Deciding if the graph is connected or disconnected is one of the first things one should look at. A regular lattice is clearly always connected. On the other hand that is not the case for a random network as we need a minimum of $(N-1)$ edges to connect N vertices (for instance in a one-dimensional lattice or in a star formation with one vertex at the centre and all others at the edge). Studying the connectedness properties of the random graphs (for large N) was part of the seminal work done by the mathematicians Erdős and Rényi³ in the late 1950s and early 1960s [13]. Their work prompted the modern era of graph theory development amongst mathematicians, as a look at one of the standard texts by Bollobás [12] shows. What Erdős and Rényi, showed was that a random graph is likely to be connected only if at least $N \ln(N)/2$ edges are present for large N . Put another way, if one increases the number of edges added randomly to a graph, then there is a sudden change in the connectedness property of the network. This ought to make the ears of a physicist prick up as it sounds like a phase transition. Indeed the analogy can be carried through much further with several characteristics of random graphs showing sudden changes as we increase the number of edges, and one can use ideas used in physics, such as percolation theory, to study these problems.

Let us continue to look at the basic ways of characterizing networks. The paths between vertices along the edges of the network have a natural length, simply the number of edges

traversed. Thus a natural measure of the *distance* between any two vertices is the length of the *shortest* path between the two vertices. If the graph is connected (or we restrict ourselves to a connected piece or *component*) then we can define an average length scale L , that is the average of the shortest distance between all $N(N-1)/2$ vertex pairs of the network. A related idea is the *diameter* of the graph which is the largest distance between any two vertices in the graph.

For a regular lattice this network defined distance between two points is called the *Manhattan distance*, as we imagine that this is the actual distance one has to travel between two places in a city with a rigid grid street pattern. For such lattices embedded in real space there is also the usual Euclidean distance between points $(0,0,\dots,0)$ and (x_1, x_2, \dots, x_D) , that is $(x_1^2 + x_2^2 + \dots + x_D^2)^{1/2}$. However, again we emphasize that this is a special property of a lattice, namely that it can be visualized as living in some D -dimensional space. It is well to note though that, as the tourist in New York knows and the example of the radio network showed, the network-defined distance and *not* the usual Euclidean one can be the relevant measure, even if the lattice network is physically embedded in real space. For such lattice problems intuition based on Euclidean distance experience can be misleading (see for example [14]) and this again reminds us that we must leave behind any pictures of our networks being embedded in some real physical space.

So, how do a regular lattice and a random graph of the same number of vertices and edges compare? Roughly speaking the size⁴ of a D -dimensional lattice will grow as $N^{1/D}$ while the size of the random graph grows much slower as $\ln(N)$. In a large random network every node is much closer

³Solomonoff and Rapoport [10, 11] had introduced random graphs and proved one classic result in 1951 but this did not seem to be well known in the mathematical community.

⁴As measured by diameter, average distances or other similar distance measures.

to all the other nodes than in a comparable sized lattice. Thus, for the earlier example of the radio communication network a random network is better than a regular lattice.

3. Local vertex properties

So far ideas such as connectedness, average distance and diameter are reflecting the properties of the network on a large scale. For many problems the local neighbourhood of vertices may play a vital role. This is familiar to physicists where the coordination number of a lattice—the number of nearest neighbours—is important in many problems and provides a simple way of highlighting the differences between say a triangular and a square lattice. In graph theory neighbours are defined only by the edges of a vertex so the number of nearest neighbours is the number of edges attached to a vertex, and this is called the *degree*, k , of a vertex. The average degree K is simply $K = 2E/N$ where E and N are the numbers of edges and vertices respectively. Thus for a sparse graph this number should be of order one, or at least K must not grow as fast as N as we increase the size of a network. For simple lattices all vertices have the same degree: in D dimensions it is $K = D(D + 1)$ for a hyper-tetrahedral (equilateral triangle faced) lattice, $2D$ for a hyper-cubic (square faced) lattice. However most other networks, such as the random graph, have vertices of a wider variety of degrees. This is best summarized by giving the *degree distribution* $n(k)$, the number of vertices in the network with degree k . This is a delta function for a lattice, but for many other graphs it is a distribution with a tail.

For the random graph it is simple to show that the degree distribution is binomial with $n_{\text{rand}}(k) = Np^k(1-p)^{(N-1-k)}\binom{N}{k}$ where p is the probability that any given pair of vertices is connected. On average there are a total of $E = pN(N-1)/2$ edges with most vertices having a degree within $3K^{1/2}$ of the average degree $K = p(N-1)$. For large N and fixed K , this is to a good approximation a Poisson distribution $n(k) \approx NK^k \exp(-K)/k!$ whose tail falls off slightly faster than an exponential, roughly as $n(k) \propto \exp(-k \ln(k)) = k^{-k}$. The fast fall-off means that there are essentially no vertices with degree bigger than k_c where $n(k_c) = 1$. Solving we find that⁵ $k_c \ln(k_c) \approx \ln(N)$ which is in most practical cases close to the most likely degree for a vertex K . For instance a random graph with $K = 4$ and $N = 10^6$ will probably have no vertex with degree larger than 17, much smaller than the potential $(10^6 - 1)$ edges available. This is clear from the distribution shown in figure 5.

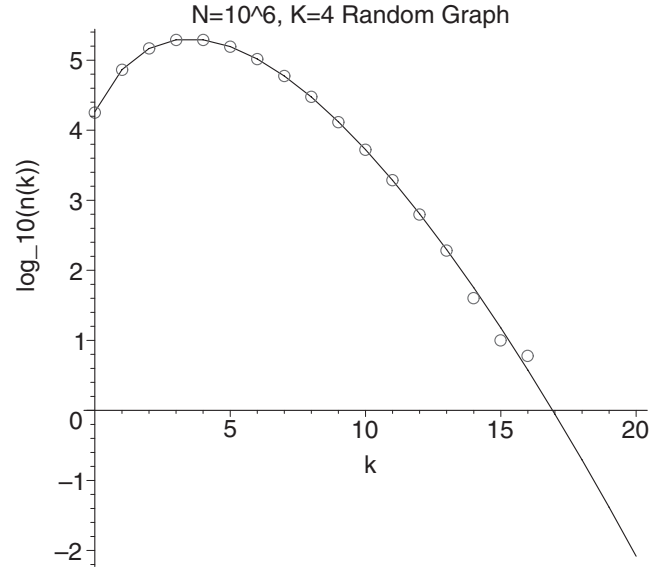


Figure 5. Plot of degree distribution $\log_{10}(n(k))$ against degree k for random networks of with $> N = 10^6$ vertices and average degree $K = 4$. The line is the most likely distribution for a random network, a binomial distribution. Circles are for a single example of a random network generated by computer and in this case 6 vertices had the largest degree of 16. The probability of finding a node with degree above twenty falls off as $\exp(-k \ln(k))$.

3.1. Clustering

While the degree distributions show clear differences between lattice and random networks, there is another measure of their differences which captures some of the order present in the neighbourhood of a lattice vertex but missing in a random graph. This is the *clustering coefficient* c . In the simplest form we pick a vertex and look at all of its neighbours, the k vertices connected to it by an edge of the network. We then see what fraction of the $k(k-1)/2$ possible neighbour-neighbour edges are present in the graph and this is c , see figure 6. On a regular triangular lattice in 2D one would have $c = 6/15 = 0.4$, reflecting its close knit local structure. Of course a square lattice has $c = 0$ and to see its tight local structure we would have to generalize our definition to c_2 involving neighbours and next-to-nearest neighbours⁶. However, in most networks the clustering coefficient c varies from vertex to vertex and an average is usually quoted⁷.

⁶We could look at the set N_2 of all nearest neighbours and next-to-nearest neighbours, and define c_2 to be the fraction of possible edges between these vertices which are actually present in the graph. So on a square lattice in two-dimensions every vertex has 12 vertices in its set N_2 with 55 possible link edges to them of which only 12 are present so $c_2 = 12/55 \approx 0.22$.

⁷As the clustering coefficient c is not defined for vertices with a degree less than two, the averages quoted in this article are taken only over vertices where c is defined. There is a second weighted average which is useful and often encountered in the literature but unfortunately the two definitions are not always clearly distinguished [15].

⁵A more accurate formula is $k_c(\ln(k_c) - 1) = \ln(N) - K$.




Figure 6. The simplest cluster coefficient illustrated for the diamond vertex v_2 . It has four nearest neighbours, v_1 , v_3 , v_4 and v_5 . Of the 6 possible edges between the neighbours of v_2 , only one, the v_3-v_4 edge is present. The other potential neighbour-neighbour edges, indicated by dashed lines, are not present. Thus the clustering coefficient of v_2 is $c = 1/6$. The vertex v_6 plays no role in this calculation. The clustering coefficients of the vertices v_3 and v_4 are $c = 1$ while for v_5 it is $c = 0$. The clustering coefficient is not defined for vertices with only one neighbour such as v_1 and v_6 here.

In some problems local communication between neighbouring vertices is essential. In lattice Monte Carlo simulations, the algorithm requires a large amount of information to be shared between neighbouring lattice sites, reflecting the local spatial (and possibly temporal) nature of interactions in many problems, such as atomic spins of a crystalline material. Thus in a parallel computer built for this problem, the communication network linking the processors should have a large amount of local structure, and indeed a lattice configuration is a common solution. On the other hand, the internet today, and computing networks such as the Grid, are designed for a great range of tasks which require different computing/data centres (the vertices) to be able to communicate easily with most other sites at different times. Local structure is not paramount. The human brain faces similar issues. The translation of the sound of a spoken command into an action involves processing in several different areas of the brain, each area specialized and involving many neurons, yet each area must communicate with many others.

The random network lies at the other extreme from a lattice and has the minimum local structure. In a random network there is no preference for neighbours to connect to each other rather than anyone else. In other words, all edges are treated equally and they are present with probability p and not present with probability $(1-p)$. In a lattice, neighbours of neighbours of some chosen vertex i

cannot be very far away in terms of Euclidean distance⁸, no more than two lattice spacings. This means that for a lattice we are guaranteed that the neighbours of neighbours are also going to be close to the original vertex i in a Euclidean distance sense, and so may well be connected to them. They will have a far higher probability than a random chance $2E/(N(N-1))$ of being connected. In a lattice, vertices are clustering together in a network sense. On the other hand, in a random graph there is no reason why the neighbour of a neighbour of the original vertex i has any other relation to i . There is no space we can use to embed our network in, no Euclidean distance argument to use. Thus we should expect that a random network should have a much lower clustering coefficient than a similar sized lattice, reflecting the fact that we can always rearrange a lattice network so it looks like the usual regular arrangement in real space so all neighbours are close to each other in all senses⁹. Indeed a quick calculation shows that a random graph of N vertices, has $N(N-1)/2$ possible edges but if only E edges are present, then on average a vertex will have a clustering coefficient of $c = 2E/(N(N-1)) = K/(N-1) = p$, just the probability p that any one potential edge in a random graph is present. For a sparse random graph of a reasonable size, this is much less than 1. Indeed, while clustering coefficients are constant for a lattice as it is made larger, for a random graph of the same number of vertices, N and the same number of edges $E = KN/2$ as the lattice, since K is fixed for a lattice, the clustering coefficients drop as $1/N$.

3.2. Small worlds

We have seen that large sparse random graphs have as little local network structure as is possible for a network since all edges are treated equally so there is no special treatment for neighbours. However this also means that the distance is small. All vertices are treated equally so each new vertex you visit on a path is likely to throw open K new vertices unreachable until now (at least for the first few steps on a route). A random graph has a small diameter. On the other hand, the way a lattice of the same size uses its edges to provide local structure, means that it takes a long time to reach an arbitrary vertex, and so a lattice has a large size in network terms. These differences are exacerbated as we consider networks with more and more vertices. On the lattice, the average degree K is fixed, as the local structure is constant by definition, but the size of the network grows as a

⁸We will assume for simplicity that our lattice exists in real space as a spatial lattice.

⁹It does not matter if our lattice network really exists as a lattice in a real space. The network properties are the same whether we look at the right hand or left hand version of figure 2. Its just much easier to visualize the local network properties of a lattice if we exploit the picture on the left.




Figure 7. Evolution of a 20 node, 40 edge network through rewiring. On the left, network A is a regular one-dimensional lattice with neighbours and next-to-nearest neighbours connected by edges. In the middle, B is the same graph with 5 randomly chosen edges reassigned, and finally on the right, C is the same graph after 200 such rewirings. The order of the vertices around the circle is the same in all cases and no more than one edge between any vertex was allowed. From left to right the average distances are $L = 2.89, 2.35, 2.21$, average clustering coefficients are $c = 0.5, 0.40, 0.23$ and diameters are 5, 4, 5, respectively.

fractional power $N^{1/D}$ where N is the number of vertices and D is the dimension. In the corresponding random network of the same number of edges and vertices (fixed K and growing N) the probability $p = K/(N-1)$ of any one link being present is dropping so the clustering coefficients tend to zero. However, the distances rise only as $\ln(N)$, much slower than any lattice.

This poses a question. Is there a sparse network which for some given number of vertices N and edges E has the local structure of the lattice but the small network size of the random graph? This is particularly important when we think about human networks. In 1967 Milgram [16] reported on an experiment in which he asked people in Omaha, Nebraska and Wichita, Kansas to send packets to one of two people in Cambridge, Massachusetts specified by name, profession and rough location only. However, packets were to be passed only between people who know each other by first name. In this way, Milgram hoped to map out the social network of close friends. The result reported was that if the packets arrived at the correct person, they had been through about five intermediaries. This is much smaller than we might have guessed given the physical distance between the original senders and the final recipients, who also presumably had no direct social or other contacts. This is the idea of the *small world*, namely that though we may think of ourselves as living amongst a small group of friends and colleagues, we all have a few contacts outside any circle and through these friends, and their friends of friends we are no more than a few handshakes from any person in the world. In terms of networks, vertices are people and edges represent links between people who are on first name terms. Milgram's experiment showed that the

average distance between sources and targets, if packets were delivered successfully, was no more than six. As Guare [17] puts it in his 1990 play

'I read somewhere that everybody on this planet is separated by only six other people. Six degrees of separation. Between us and everybody else on this planet. The president of the United States. A gondolier in Venice.' ...

'It's not just the big names. It's anyone. A native in a rain forest. A Tierra del Fuegan. An Eskimo. I am bound to everyone on this planet by a trail of six people. It's a profound thought.' ...

'How every person is a new door, opening to other worlds. Six degrees of separation between me and everyone else on this planet.'

This 'six degrees of separation' between people reported by Milgram has captured people's imagination and has entered modern cultural mythology¹⁰.

In fact this type of behaviour is common in many areas of human interaction so references to it are not uncommon, and certainly predate these examples¹¹. I remember some

¹⁰So it is intriguing to learn that these small world ideas championed by Milgram and certainly of great significance may not actually be well supported by all the data he collected, both that reported in the scientific literature and unreported results in the Yale archives of his work. Watts [18] gives an excellent overview of the work of Kleinfeld [19] who reviewed the results Milgram actually obtained in his various small world studies.

¹¹Barabási [20] mentions an obscure short story 'Láncszemek' (Chains), paradoxically by a famous Hungarian author Frigyes Karinthy published in 1929 where it was suggested that it took five acquaintances to reach anyone.

excited students showing me the ‘Kevin Bacon game’ and its web site [21] for the first time. Name an actor, and the web site will tell you their ‘Bacon number’, the number of steps it takes to reach Kevin Bacon via actors who are paired if they appeared in the same film. For instance Charlie Chaplin has a Bacon number of 3 since he appeared in ‘Screen Snapshots: Spike Jones in Hollywood’ (1953) with Douglas Fairbanks Jr, who in turn was in ‘Hollywood Uncensored’ (1987) with Eli Wallach and Mr Wallach was in ‘Mystic River’ (2003) with Kevin Bacon. The actors are vertices, edges represent a film in which both the actors associated with ends of the edge appeared. This amusement shows that paths between actors are surprisingly short, around 3.7 on average (in 1999 [22]) with lower numbers for actors such as Kevin Bacon who at the time of writing was only 2.944 from another actor on average¹². However, long before computers and the internet enabled this game to be played with ease, the mathematicians had the same concept of an Erdős number, the same Erdős that produced the seminal random graph papers. The vertices were mathematicians who were linked if they had co-authored a paper¹³. Erdős co-authored with a large number of people, a result of a fabled itinerant lifestyle, an innate mathematical ability of his own and apparently an ability to encourage and stimulate the work of others. The mathematician’s game is a tribute to Erdős [24].

The underlying idea in all these examples is that these networks reflect human social interactions. They have a lot of local structure, many of our friends are our friends’ friends too, many actors appear in the same type of films which therefore draw on a small pool of actors active in that genre at the time, academics often collaborate several times with the same people. In terms of networks, they have a relatively large clustering coefficient. How large is ‘large’? Well the best way to put it is that it is much larger than a random graph of an equivalent size. On the other hand, the distances across the network are small, comparable with those obtainable from a random graph and much smaller than any regular network. This then is the definition of a small world network¹⁴, namely $c_{sw} \gg c_{random}$, $L_{sw} \sim$

L_{random} . This is all very well, but what type of model gives us a small world?

This is where the seminal paper of Watts and Strogatz of 1998 comes in [23]. Their idea was to start from a lattice, in their case a one-dimensional ring with neighbours and next-to-nearest neighbours linked, so the average degree was $K = 4$. You then look at each of the $E = 2N$ edges in turn and with probability p , you *rewire* this edge, that is you choose two new vertices at random and place the edge between them¹⁵. This is illustrated in figure 7 in what is perhaps the iconic image of this topic in its modern era. If $p = 1$ every edge is placed at random and we have a random network of the same number of vertices and edges, short distances but little local structure. If $p = 0$ we have the original lattice, high on local structure but with large distances. For intermediate p we might expect a hybrid of these features with clustering coefficient and distance to drop together as we increase p . In fact Watts and Strogatz highlighted the fact that as p was increased, the distances dropped rapidly down to random network levels, while the clustering coefficient only came down slowly at first (linearly for small p). Thus for a small range of p values in their model, they were producing a small world network retaining most of the local structure of the lattice but having the short distance characteristic of the random

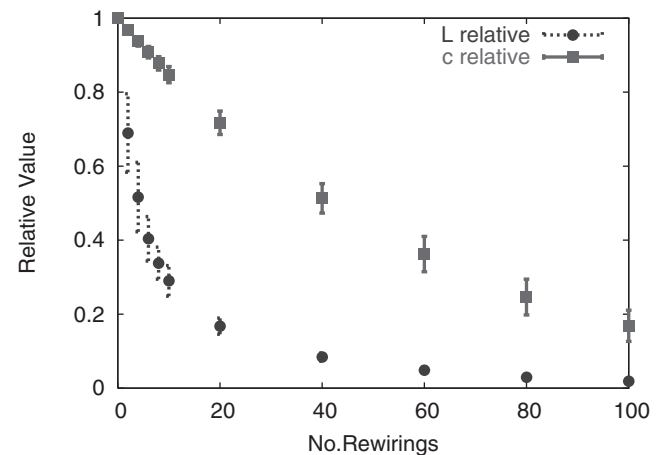


Figure 8. Plot of c_{rel} (top square points) and L_{rel} (bottom circle points) against the number of rewirings (proportional to p for small p) in a WS rewiring scheme for 100 runs starting from a 100 vertex $K = 4$ one-dimensional lattice. The value $c_{rel} = (c - c_{rnd}) / (c_{lat} - c_{rnd})$ is the clustering coefficient scaled so that it is 1.0 for the lattice (left-hand side) and 0.0 for the random graph (right-hand side). L_{rel} is scaled in a similar manner.

¹²The web site states that on 29th April, 2003 Rod Steiger had the shortest average distance, average Bacon number of 2.652, so disproving the theory behind the student bar game that Kevin Bacon was the centre of the Hollywood universe. Mr Steiger was just ahead of Christopher Lee at 2.660 while Mr Bacon’s average was 2.941 putting him only 1222nd on the list on that date. The network is constantly growing and at the time of writing, February 2004, the largest connected component had 6.2×10^5 vertices (actors) with an average degree of $K = 2.94$.

¹³Both the academic co-authorship and Kevin Bacon game examples can be played in a different way. We could have the films (academic papers) as vertices with the actors (academics) as links. Get from ‘Ben Hur’ (1926) to ‘Ben Hur’ (1959) (from ‘On Random Graphs I’ [13] to ‘Collective dynamics of ‘small-world’ networks’ [23]) in as few steps as possible. This illustrates the way that the same data can often be represented as a network in many different ways.

¹⁴Well, except that some people drop the clustering part of the definition. You just cannot win with the nomenclature in this field.

¹⁵In fact their algorithm was slightly different in implementation and there are many slightly different algorithms which achieve the same result. The key idea is adding a few random links to a regular lattice gives a small-world network while adding more random links eventually brings you to a random network.

graph. Results for an $N = 100$, $K = 4$ case are shown in figure 8. These are easy to understand qualitatively. Rewiring one or two edges will only lower the clustering coefficient of a few vertices. For the first few rewirings then, the average c is going to come down in proportion to the number of rewirings. On the other hand, while removing an edge makes little difference to most paths between vertices, putting it back is likely to create a shortcut between distant vertices on the lattice. Many shortest paths between vertices will now be shorter as many paths can take this one shortcut. The average distance drops dramatically for the first few shortcuts.

As we have noted, the idea of a small world network was not new. The WS (Watts and Strogatz) model is just one exemplary model of a small world but it offered the prospect of simple models for new types of network and a systematic way of characterizing them. This, when coupled to modern desktop computing power, allowed any scientist to make a systematic and statistical study of networks. As in the case of the random networks, analytical results are hard to come by so the ability to do numerical experiments is central to this field. The same computing power, along with the ability to gather and exchange large amounts of data, also enabled the analysis of real data along the lines suggested by the work of Watts and Strogatz. Of course, others, notably the social scientists, had been analysing such data using computers and the language and techniques of graph theory for some time. However, Watts and Strogatz opened the doors for physicists to join in as the explosion of preprints on networks placed on the archive cond-mat shows. In 2003 the number of papers containing a word starting with 'network' in their title was 730% higher than in 1997, see figure 1.

3.3. Hubs

Not surprisingly a random solution is not a good solution for the radio communication problem suggested above. One solution which can connect more sites for the same number of edges is where one site is connected to every other site, all of which have only that one edge. What we then have is a single dominant hub in our network and all paths in the network will go through this hub. This means that all vertices are relatively close to each other in terms of the natural network distance, as the path from each vertex to the hub is relatively short and then again from hub to any destination vertex takes only a few edges. While short distances are a feature of random networks and small world networks, this is not the whole story. The random network and the WS small world network do not have great big hubs. There the exponential nature of their degree distributions $n(k)$ for larger k means that the vertices with the most connections are of a relatively small

size in a sparse network. As noted above and shown in figure 5, in a random network of $N = 10^6$ vertices and average degree $K = 4$ there is unlikely to be a vertex with degree higher than 17. Thus these types of network do not allow for very big hubs, such as we might expect to see in some real world situations or with the star formation solution to the radio network problem. Of course, there are practical limitations on hubs in many problems which prevents such extreme hubs, as the centre of a star formation, appearing. For instance in communication networks dealing with too many channels simultaneously at any one point is likely to be impossible or unrealistically expensive. Still, it seems that for some problems large hubs are preferable and it should not be too surprising to find that there are many studies which seem to see many more large hubs than one would expect from random or WS small world networks.

In order to have hubs, we need a degree distribution $n(k)$ with a much longer tail than the exponential ones of random and WS type networks. One natural candidate is a power law, $n(k) \propto k^{-\gamma}$. It is a natural candidate because it has a very long history in a wide range of subjects outside physics, as well as playing a central role in several areas of physics.

For instance, Lotka [25] in 1926 claimed citations of scientific papers followed a power law. In this case the papers are vertices and citations from one paper to another form edges. This is a good example of where the internet communication and desktop computer power has revolutionized work on networks. Redner [26] studied this network in 1998 using the 1991–1997 citations of 1991 papers listed in the ISI (Institute for Scientific Information) database for his largest example. This gave him a directed network of about $N = 7.8 \times 10^5$ vertices with average degree around $K = 8.6$. A much longer tail than a simple exponential of this size and far more hubs are present. A $\log(n(k))$ versus $\log(k)$ plot reveals a linear tail for the data with a slope of about $\gamma = 3$ ¹⁶.

Networks with power law degree distributions for large degree are known as *scale free networks* because ratios such as $n(2k)/n(k)$ are constant, whatever the degree scale k is chosen to make the comparison¹⁷.

To see that such a distribution has hubs, vertices with much higher numbers of edges than in a random graph or in a WS model, consider the following example. Consider a network with $N = 10^6$ vertices and an average degree of $K = 4$ with vertices of all degrees present (so the minimum

¹⁶As I have defined small world graphs in terms of clustering and distance measures relative to a similar sized random graph, there is no clear statement one can make about the degree distributions of small world networks.

¹⁷Always provided that we are looking at vertices with large enough degree where $n(k)$ is a power law. The degree k_* where power law behaviour starts does in fact provide one scale for the problem.

degree is 1) whose degree distribution is of the form $(k + k_s)^{-3.0}$ ($k_s \approx 2.91$) and so for large degrees it's just a simple inverse cubic¹⁸. It is easy to show that for every value between 1 and about $k_c \approx 284$ there is at least one vertex of that degree ($n(k) > n(k_c) = 1$)—a continuous spectrum of degree. Once the likely number of vertices with a given degree k falls below one, $k > k_c$, then we find that in a real network, where vertices have integer numbers of connections¹⁹, the $n(k) = 0$ for many $k > k_c$ and we are at the end of the ‘continuous’ part of the degree distribution. However, unlike the random graph, the probability of finding vertices with degree bigger than k_c is not falling off very fast and there is a non-trivial chance of finding the odd vertex with a large number of edges attached. In this

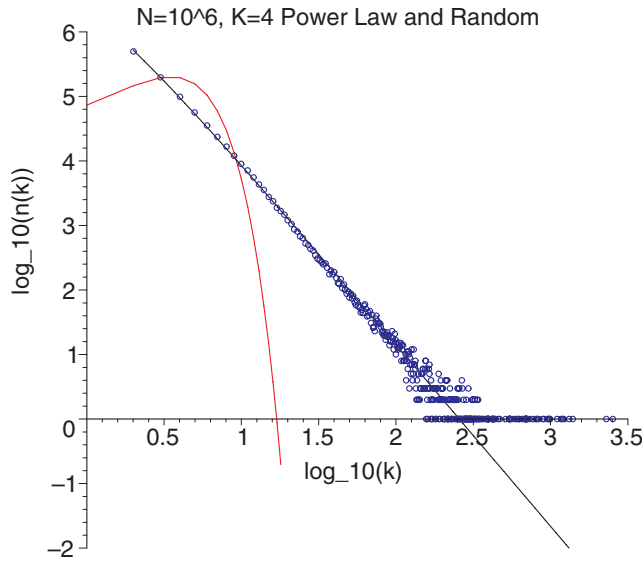


Figure 9. Log–log plot of degree distributions for a network with $N = 10^6$ vertices and average degree $K = 4$. Data from a model of a scale free network (blue circles, walk algorithm) compared against a power law with $n \propto (k + 0.42)^{-2.82}$ (black straight line) (both with minimum degree $k_0 = 2$). Note how the power law crosses $n(k_c) = 1$ at about $k_c \approx 257$ ($\log_{10}(k_c) \approx 2.41$) and this is roughly where the data for the real network is no longer continuous. Even if we used tricks to extract data for higher k the highest data point is likely to be around $k_1 \sim 4800$, though here it is at 2519, giving us only two and a half decades to make the linear fit. Note that having a minimum degree of 2 rather than 1 shifts the power law to the right compared to examples where the minimum degree is 1. Also shown is the binomial distribution of a random network of the same number of vertices and edges (red curve) whose largest hub is at $k \approx 17$.

¹⁸If we specify N , K , γ and $k \geq k_0 = 1$ then we need one last free parameter, say k_s , to allow a fit. In this case we must have $k_c \approx 2.91$. In the following we analyse this distribution and its various sums over integers by noting its simple relation to the Hurwitz–Riemann zeta function whose properties are well known.

¹⁹The formulae for degree distributions $n(k)$ imply that some statistical average over many realizations of networks is taken.

example we expect to find about 143 of these large vertices with degrees between about $k_c \approx 284$ and $k_1 \approx 1447$ where we can show that there is unlikely to be a vertex with degree larger than $k_1 \approx 1447$ ($\sum_{k=k_1}^{\infty} n(k) = 1$). Plots for a similar distribution from a real model are shown in figure 9. These numbers, k_c , k_1 , are much larger than those for a random network of same size where both are approximately the same and are around 17. In fact for a power law distribution these numbers are scaling as fractional powers²⁰ of N while the random graph cut-offs are roughly equal and scale roughly as $\ln(N)$.

If the citation networks of Lotka and Redner were the only reports of a power law distribution in the literature, then perhaps we could merely dismiss it as a special and rare case. However Newman [27, 28] and others have reported similar patterns in scientific collaboration in different academic fields, the modern version of the Erdős number discussed above. Albert *et al.* [2, 29] showed that the network of links (the edges) between web pages (the vertices) seems to be a scale free network. Three Faloutsos brothers showed that the computers and routers that form the internet itself form a scale-free network [30] ($\gamma \approx 2.2$). Patterns in long distance phone calls [31, 32] ($\gamma \approx 2.1$), the network of film actors used for the Kevin Bacon game ($\gamma \approx 2.3$) [33], the relationships of words in English texts [3] ($\gamma \approx 2.7$) and many other examples have been given in the recent literature. They all have distance scales comparable to a random graph but much larger clustering coefficients, i.e. they are also small world networks. Indeed while not all networks are scale free, many in the literature are reported to be just that.

A power law distribution is something that should again make all physicists sit up. At or near a second-order phase transition, many quantities follow a power law. For instance the magnetization of a real magnet near the critical point is proportional to $(T_c - T)^\beta$, where T is the temperature and T_c is the critical temperature. Physically the scale invariance reflects the fact that near a critical point the relevant correlation length becomes infinite, there is no relevant scale left. Further, this means the small scale details of the material become irrelevant so one has universality, many materials with different short scale interactions have the same behaviour near the critical point. This insight led to powerful new ways of solving problems near critical points. Subsequently, whenever a power law is seen one hopes that this reflects the emergence of a similar universality and so the hope is that if a simple model can capture the essential physics, it will be guaranteed to give the correct results for the real world examples, just as simple Ising models are good models for the critical behaviour of real materials.

²⁰With the end of the continuum at $n(k_c) = 1/N$ and the largest degree around k_1 where $\sum_{k=k_1}^{N-1} n(k) = 1$, we find that $k_c \propto N^{1/\gamma}$, $k_1 \propto N^{1/(\gamma-1)}$.

Power laws are also central to more recent ideas. For instance the length of a coastline l measured with a ruler of size k is also a power law, $l \propto k^{-\gamma}$. In this case it reflects the fact that the coastline is not a simple line of dimension one but a *fractal* with a fractal dimension γ . In the 1980s and 1990s, power laws were at the centre of interest in various complex systems. The distribution of the number of earthquakes n of a certain size (equivalent to k) is a power law over several decades (Gutenberg–Richter law) even though describing the properties of the Earth’s plates, their response to forces, their interactions with the core, appears to be far too complicated to give such a simple result. A clue comes from the existence of several models where the microscopic rules are very simple, even if the equilibrium solution may be hard to find analytically. Numerical or even real experiments can be performed to show critical behaviour with power laws. For instance in the sand pile models where ‘sand’ grains are dropped regularly onto a pile, there is a rule that the sand is stable unless the height difference between its neighbours exceeds some critical point. This produces a number of avalanches (n) which is a power law of the size of the avalanches (k). One idea seen in some models is that many systems may actually prefer to lie on a critical point, one can alter the initial conditions or the microscopic rules but they always give a power law if one waits for some time. This is called SOC—*self-organized criticality* [34, 35].

Intriguingly, many power laws occur in areas outside physics, and several famous ones predate physicists

interests’ in networks, SOC or critical phenomena by a considerable time. We have already noted Lotka’s interest in citations [25]. He observed that the number of authors contributing n publications to a bibliography is often an inverse square law, but did not look at this in the context of a network (papers as vertices, citations as links).

Many power laws in social science are said to be examples of Zipf’s law (see figure 10). George Kingsley Zipf [37] noted that for many quantities if you rank them in order of size, largest given rank one, second largest rank two, etc. and plot their size against their rank, then we get a power law. One example he gave was of the frequency of English words and he suggested the frequency of a word was approximately inversely proportional to the rank. The universality of this idea forms the basis of many data compression algorithms. His second famous example is the sizes of cities and again suggested that city size was inversely proportional to rank. We can turn our degree distributions into the language of Zipf’s law as the rank γ of a vertex of a network, in terms of its degree k , its ‘size’, is given by $r = \sum_k n(k)$. Zipf’s law then consists of plotting k against r . One can quickly see that $k \propto r^{-\nu}$ with $\nu = 1/(\gamma - 1)$ for a power law distribution $n(k) \propto k^{-\gamma}$.

The characteristic of all these power laws can be summarized in terms of an older ‘law’, Pareto’s law. Pareto was an economist working in the late 1800s who noted that 80% of the land in Italy was owned by 20% of the population. Such an 80:20 rule applies whenever we have a distribution with a significant tail such as a power law

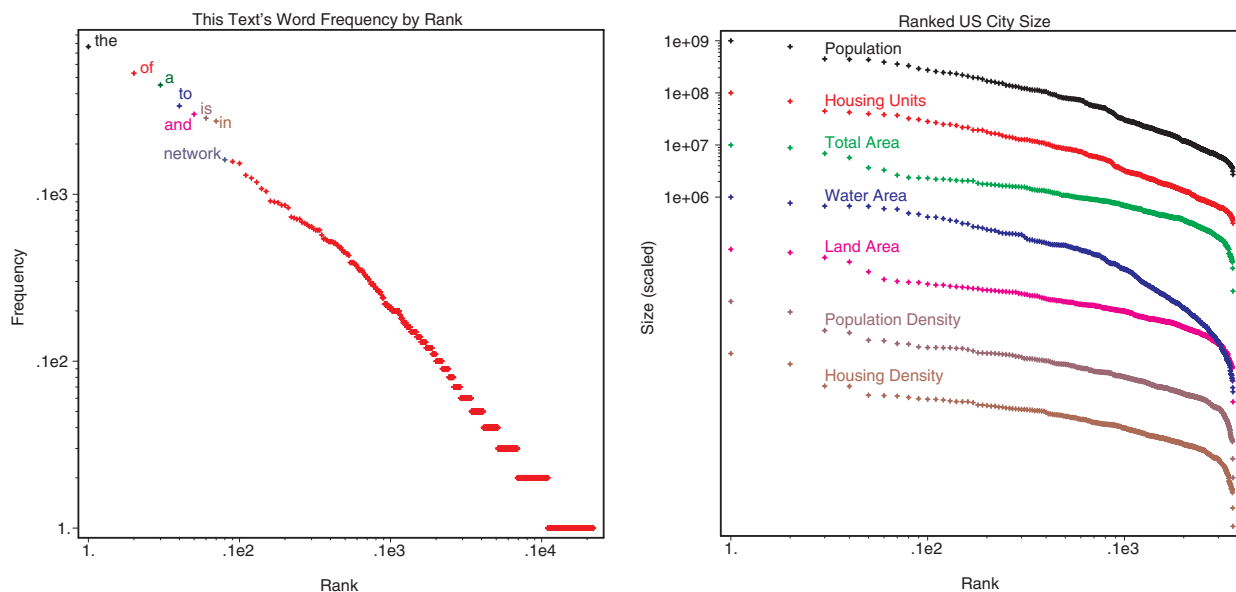


Figure 10. Two classic Zipf plots. On the left a log–log plot of word frequency against rank for this text—‘the’ is the most common word, ‘network’ is eighth. On the right, a log–log plot of the size of US cities against rank, all data scaled relative to the largest and by a power of ten relative to the next curve (for visualization purposes). The data is for ‘Metropolitan Areas’ in the USA in 2000 and taken from the US Census Bureau web site [36]. Size is given seven different ways in terms of population, housing units, total area, water area, land area, population density and house unit density. Most show some evidence of a simple power law.

distribution. For instance only a few English words are needed to write most of a book, Lotka's study means only 40% of authors have more than one publication in a bibliography.

Finally, one of nicest examples of scaling laws are those relating to metabolic function observed in biological systems. For instance the lifespan of an organism scales as its mass to the power one quarter and this relationship holds over 21 orders of magnitude, from microbes to whales. It is particularly relevant as the very ideas of scaling used so successfully in physics to explain critical phenomena have recently been used by West *et al.* [38, 39] to provide a simple explanation for these previously mysterious biological scaling laws.

So is it possible that when we see scale-free networks, networks with a power law degree distribution, it reflects some simple fundamental law that could be described by the network equivalent of an Ising model? Are most of the details in network creation irrelevant and we need only focus on some crucial ingredient, something which means many networks will always eventually organize themselves in a special way and the scale-free degree distribution is flagging this? Are networks fractal in some sense? Not all the power laws above, such as the city data, have an obvious network basis but perhaps we would gain great insight if we could see their power laws as a result of some scale-free network structure which we have missed to date. Many of these power laws appear in areas well outside traditional physics, social networks with their six degrees of separation and so forth. Can the language of networks provide us with the first clues of some rigorous theory of some aspects of human social behaviour—a realization of Asimov's psychohistory [40]? Or is this destined to remain science fiction?

3.4. A word of caution

These ideas about power laws, these questions, these speculative links certainly provide one explanation for the excitement in the study of networks in recent years. However, a word of caution. Modern computers and the internet have allowed researchers to create and analyse large databases across a wide range of topics. It may seem that a data set of 10^6 vertices is large (and I can think of only three examples which are much bigger). However the degree distribution is much shorter than 10^6 . Suppose our system gives a network of $N = 10^6$ vertices, average degree $K = 4$ and a minimum connectivity of $k_0 = 2$ and a $(k + k_s)^{-\gamma}$ form for the degree distribution with power $\gamma = 3$ and scale $k_s = 0.87$ so it is effectively a pure power law for almost all k . One typical example of such a network will have a non-zero degree distribution $n(k)$ only for k below $k_c \sim 226$ (where $n(k_c) = 1$). Above this value there are often no vertices with that particular degree k , $n(k) = 0$,

with only the occasional $n(k > k_c) = 1$. So plotting this directly on a log–log plot we have to ignore many points where $n(k) = 0$ while the few remaining ones are all the same value, one. For similar reasons the region just below k_c also shows large fractional fluctuations. The data from a computer model of a similar scale-free network in figure 11 shows this clearly. Thus even in the best case where there is a power law for all values of k (scale free networks need be power law only for large k) we have only two decades of linear behaviour in our log–log plot. Anyhow, in most cases the power law form starts at some degree higher than one.

There are various tricks which improve the situation a little. One is to bin the data in the large degree $k \gtrsim k_c$ region as in figure 11. That is if we count η vertices of degree between k_{low} to k_{high} , one 'bin', we assign a single data point of value $n(k_b) = \eta / (k_{\text{high}} - k_{\text{low}})$ at²¹ $k_b = (k_{\text{low}} + k_{\text{high}})/2$. If we choose these bins so that $k_{\text{high}}/k_{\text{low}}$ is a constant, we get a series of equally spaced points on a log–log plot. By averaging over a range, we have fewer data points but they are no longer limited by the discrete values 1 and 0 and they allow us to fit in this region. However, with the largest degree around $k_1 \sim 2400$, $\log_{10}(2400) = 3.4$, we will not add more than a decade to our linear region. As an alternative, one can try to plot a Zipf-like size versus rank distribution, see figure 11. However, both binning and ranking introduce correlations between data points so the statistical analysis is not so simple.

It is not surprising then that many authors do not put an error on the power law fitted to the data, nor can one always exclude other types of fit to the data, stretched exponentials $k^{-\gamma} \exp \{-(k/k_s)^\beta\}$ for instance. It would be nice to get more data but it is only with modern computers that networks of a million vertices can be easily found and, even then, I know of only two studies of the web and one of phone calls where the number of vertices is larger (though still no more than 10^8). Many data sets are smaller and are not going to get significantly bigger. The history of power laws in other areas of modern physics tells us that they are so intriguing that we sometimes rush to see power laws in everything and this is premature in some cases. Another consequence is that we are not going to see power laws with large coefficients [41]. If we require at least two decades of linear data and assume a simple power law for the whole range of $k \geq 1$ then we require that $k_c = 100$, where $n(k_c) = 1$, and we need $N - 100^7$ vertices. Thus with $N = 10^6$ we should be able to see scale free behaviour in networks if the power γ is less than 3. So while the ranking and logarithmic binning can help improve the data, the power law section of accurate data is rarely going to be even this

²¹The optimal position chosen depends on the nature of the curve, $k_b = (k_{\text{low}}k_{\text{high}})^{1/2}$ is another choice but in practice the two are essentially the same for $k_{\text{low}}/k_{\text{high}} \ll 1$.




Figure 11. On the left a log–log plot of the degree distribution from a theoretical model of a scale-free network with $N = 10^6$ vertices, average degree $K = 4$, minimum degree 2. The best fit gives a power of 2.82. The black crosses are the raw data and the blue diamonds are the result of logarithmic binning. Note how the binning leaves the form unchanged from around $\log_{10}(k) = 1$ where it first makes a difference and has a consistent shape until $\log_{10}(k) = 3$ though the quality of any fit will be poorer at these high k . On the other hand the raw data has large fluctuations from around $\log_{10}(k) = 2$ so the binning is gaining us nearly a decade of data to fit any theoretical curve to. As an alternative, the log size (degree) versus log rank plot à la Zipf also shows a much cleaner result for the biggest 10^4 vertices with slope around -0.55 .

long and we should expect three to be an upper practical limit on powers found in real data. The worry is that with only short ranges of linear-like data, and two decades is not so long, many other forms can fit the data just as well, such as stretched exponentials. These additional forms cannot be excluded. If we are interested in the form the network would have as $N \rightarrow \infty$ then we must be very careful, as finite size effects may be difficult to detect with the data. On the other hand if we are just characterizing our network, and not ascribing some deeper meaning to an exact power law relation, then we need not be so worried. Just by looking by eye at many of the data sets, one can be convinced that the $\log(n(k))$ versus $\log(k)$ plots do have convincing linear tails and the identification of scale-free networks in a wide range of networks in the real world is not unreasonable.

4. So what?

So far we have discussed several key aspects of a network which allow us to distinguish different types of network. While classification is important, let us now try to see if there are further useful questions we can answer about networks.

4.1. Optimization

There are several types of question traditionally asked about networks. There are the ‘optimal route’ questions such as the travelling salesman problem mentioned at the start of the article. Perhaps the oldest example is the one that is seen as the start of graph theory in mathematics. In 1736, the great Swiss mathematician Euler proved that it was impossible to walk round the city of Königsberg

(modern day Kalinigrad) crossing each of its seven bridges across the River Pregel once and only once, a preoccupation of some of its citizens at the time. This is equivalent to asking if you could walk along every edge of a particular network once and only once. Critical path analysis is another well-established problem. If we use a graph to represent the different tasks in a problem, each depending in some way on earlier parts of the problem to be complete as indicated by directed links, can we find the bottle necks in the process so we can focus our resources on these and complete the job in the shortest possible time? Problems such as these have a long tradition and finding algorithms which give one good answer in a reasonable time, rather than finding the best answer at any cost, preoccupy many computer scientists. The application of such optimization problems to a wide range of problems predate the current interest in networks. For instance Davis [8] used network methods to suggest that the pivotal role of the island of Delos in ancient Aegean culture, when it is a rather small and insignificant island, was due to its critical position in the sea-borne communication networks of the era.

However, let us look at some of the questions that the recent interest in networks has provoked or revived.

4.2. Where do networks come from?

One of the first questions that comes to mind is why do different types of network appear in different contexts? How are different types of network formed? This may be related to the origin of some of the power laws found in a wide range of human experience, city sizes by rank, river size against river basin areas, the Gutenberg–Richter law of earthquake size–frequency, etc. Is there a simple guiding

principle behind the patterns in systems with such complex interactions and can this be related to some type of network? Here there is real hope. The story of the elucidation of the simple physical and biological principles which can explain the power laws seen in biology [38, 39] offers an exemplary model, at least in the way most physicists view problems.

For the lattice we have regular solids with short-range potentials and a minimum energy principle to guide us and these do not interest us here. The random graph and WS models were presented in terms of an algorithm for their creation. Thus we can see how small world networks might be created. However, are most networks going to be formed by purely random interactions (rewirings)? While there may be an element of chance about which web pages a web page author reads and therefore which links that author is likely to add to their pages, there is also a large amount of informed decision making going on. The author reads web pages related to his/her interests, not just any old page, and they link only to such pages in general. Thus the Erdős and Rényi, random graph model and the Watts and Strogatz algorithm for small worlds while useful for theoretical analysis are perhaps not good models of the processes that lead to networks in the real world.

One significant recent contribution has been to provide a model for the creation of scale-free networks. As we have seen, these have been of interest for many years, although usually interest focused on the power law and not any underlying network. One can find algorithms for creating scale-free networks e.g. [32], which involve the input of the power law form but this is not going to help us study their origin in the real world. A major insight came from Barabási *et al.* shortly after the Watts and Strogatz work on small worlds. Their study of the world wide web [2] showed a network with lots of hubs, incompatible with random graphs. Seeing that the web was growing at a considerable rate, they suggested a model that had two key ingredients, growth and what they called ‘preferential attachment’. Imagine adding one new vertex at each turn to an existing network. To this vertex we attach $K/2$ new edges with one end connected to the new vertex, the other end connected to an existing vertex chosen with probability Π from all the vertices in the existing network. If $\Pi = 1/N$ so that all existing vertices are treated equally, we end up with an ‘exponential’ network²² and there will be no hubs and no vertices with far more edges attached than the majority. To reach a scale-free model Barabási and Albert [33] suggested that these new edges attach preferentially to vertices with large degree k . A ‘rich get richer’ algorithm, echoing Pareto’s law. The simple form they took was $\Pi(k) \propto k$ where the normalization is simple to calculate, as illustrated in figure 12(a). What is remarkable is that this model gives a power law degree

distribution with $\gamma = 3$. In fact this model is a simpler version of models suggested by the Nobel prize winner Herbert Simon for the size distribution of businesses as early as 1955 [42–44] and there are several simple variations of this model, such as $\Pi(k) \propto k + c$ with c being a constant, which can give scale-free networks of any power above²³ 2. Moreover, many simple models can be solved exactly in a mean field approximation (see for example [41, 45–47]).

Unfortunately what is also remarkable is that in this type of model *only* probabilities Π linear in degree give power law networks. Anything else gives a stretched exponential, such as $n(k) \sim k^{-\gamma} \exp \{-(k/k_s)^{(1-\gamma)}\}$, or other forms as seen in explicit solutions [48]. It seems unlikely that we choose which web pages to link to with probability *precisely* proportional to the number of web pages on a site at that time. Still the Barabási and Albert model has highlighted several crucial aspects about scale-free networks and power laws. In particular it emphasizes that the power laws come from networks where there are hubs, that is vertices with far more edges attached than the vast majority and far bigger than found in simple random or WS type small world networks. The idea then is that the processes leading to the formation of a scale-free network *require* some type of ‘preferential attachment’, a preference for the most connected vertex, does make qualitative sense. If I think it is worth linking to a web site, the chances are several other people will think the same. I tend to link more to popular web sites. Perhaps one definition of what one means by ‘popular web site’ is a site to which many different people have made links. It is perfectly reasonable to think that we are more likely to link to popular web pages than obscure ones, and that popular ones are the ones with the highest number of links to them, the highest degree.

In fact the idea can be developed to produce more realistic models for scale-free networks [49–52]. Suppose I am writing a new web page and want to add a link to good pages about my favourite hobby. What I might do is use the web to find these pages. So I go to a search engine (a super hub surely?) which points me to various pages. I try a few, perhaps follow links from those to pages now distance two from the search engine vertex. What I am doing is walking along the edges of the network. It is the structure of the network which guides my search and thus I would expect that this guides the way I connect to it, the way the network grows. Of course, what happens is I am most likely to end up at sites that are linked to by many others. There are more ways of getting to them than the sites with only one or two links from the outside world. The very search I do is likely to lead me preferentially to the most connected

²²Not a random graph of Erdős–Rényi type but a network with a simple exponential fall-off for the degree distribution.

²³A pure power law with a power equal to two has formally an infinite average degree, informally too many edges. This is why models cannot reach powers of two and below.




Figure 12. Two of the algorithms for creating scale-free networks. In both cases, we add one new vertex v_7 (green diamond) and want to add one new edge between it and one of the vertices in the existing network, v_1 to v_6 . The most likely edge to be chosen is indicated with a dashed green line and is to the one with highest degree, v_2 , in both cases. In (a), we illustrate the algorithm discussed by Barabási and Albert in which we connect our new edge to an existing vertex with probability proportional to the degree of that vertex. These probabilities are given in square brackets. In (b), the walk algorithm is illustrated. We choose at a vertex from the existing network at random, say v_3 , and then start a random walk from this point. A typical example is indicated by the arrows. We stop after a certain number of steps and connect the new vertex to the vertex at the end of our walk, v_2 in this case.

sites. Indeed, if we idealize this, and suppose that we execute a random walk on the network, assuming we walk for a distance $d \gtrsim L$ the mean distance, I am likely to arrive at a vertex independent of my starting point and the probability of me finishing at a vertex with degree k will be proportional to k as there are k different ways of arriving at this vertex²⁴. Thus if we were to attach a new vertex to a vertex in an existing network where the latter is chosen after a random walk on the lattice, we will quickly generate a scale-free network. Such a ‘walk algorithm’ was used to generate the scale-free models in figures 9 and 11 and is illustrated in figure 12(b).

Growth is also not a requirement for scale-free networks, one just needs an ongoing dynamical rearrangement. One approach is to rewire as in the original WS model but now we choose how to rewire an edge in a more general way, e.g. choosing to reconnect existing edges to new vertices of degree k with probability $\Pi \propto k$ [53]. More generally we can steal ideas from statistical physics and define a Monte Carlo type algorithm in which the network emerges as an equilibrium state. The idea would be that we give the network an ‘energy’ with the Hamiltonian containing terms proportional to the number of vertices, the number of edges and other suitable terms. We would try removing or adding an edge or a vertex, and use the Monte Carlo algorithm to accept or reject such an update. The coefficients of the terms in the Hamiltonian play the roles of temperature and chemical potentials and we could imagine fixing some of these terms, the equivalent of working in different types of

ensemble. The equilibrium solution gives the number of vertices with a given connectivity rather than the number of states with a given energy and this allows us to see which type of terms in the Hamiltonian are required for different types of network. In the real world, such a Hamiltonian may arise because we have costs, each edge of the internet might be a cable which costs money to use.

4.3. *Desperately seeking ...*

In the quote from Guare’s play, I left out a couple of sentences. Of relevance here is a comment following the first part on the six degrees of separation

‘I find that A) tremendously comforting that we’re so close and B) like Chinese water torture that we’re so close. Because you have to find the right six people to make the connection.’

while the last part continues

‘But to find the right six people.’

Guare is highlighting an important aspect of Milgram’s experiment, namely how we search the network to find the shortest paths²⁵. If there were only six degrees of separation between people in Nebraska and Massachusetts, how did they find this route? There are infinitely many ways that a

²⁴Strictly speaking this is using a mean field approximation.

²⁵Again note the discussion in Watts [18] and Kleinfeld [19] about the actual data Milgram collected.

letter could have been passed along. Indeed the fact that only about 20% of the letters sent actually arrived could be in part be due to the fact that some people did not find a good route. Further, perhaps there were shorter routes but people were not able to find them. There is not much point in having a network if we cannot find our way about it. Finding good routes will be an essential part of almost any network. In anthropological terms, many cultures set up a network of contacts, for example through marriage or gift giving. One use of such a network is the right to use it to find a skill or resource you need but do not have. In terms of computing, one way of increasing computer power is to distribute the computing and storage facilities across many computers. The peer-to-peer file sharing networks can be run without central coordination but then how do you find which computer has the file you want? If one sends out a message to all your neighbours in the network, and they in turn send out to all their neighbours if they do not have it, the network will be swamped by such requests [54].

Consider for example Kleinberg [55] who considered a two-dimensional lattice to which short cuts have been added with probability proportional to $d^{-\alpha}$ where d is the Manhattan distance (network distance if there are no shortcuts). He showed that if you know the Euclidean position of the target and if the short cuts satisfied an inverse square law $\alpha = 2$ then with a simple algorithm he specified which used only knowledge available locally at each vertex²⁶, it was possible to find short paths from initial to target vertex, and it would take about $(\ln(N))^2$ time (equals the number of steps taken on the network). For any other situation, a different algorithm, different power law for the short cuts, the time taken would be considerably longer.

Kleinberg's work is a good example of what can be achieved analytically and it is relevant for many problems. However knowing that there is an underlying lattice and using the Manhattan distances is in a sense using some global information about the network. In Milgram's experiment, this might have helped to some degree, a letter from Nebraska might have been sent first to a friend in Boston, as this is physically closer to New York, who then forwarded to a friend in Massachusetts. However once in the final geographical region, another way of finding the target is needed as only rough geographical information was given. Presumably the searches in this case were then made by using intuitive ideas about the distance between different professions or something similar. For instance if the target is a medical doctor then I give the message to any friends who are medical doctors or, failing that, I give it to any one who works in the medical profession as while they are not as 'close' professionally, it is natural to think of them as 'closer' than other friends who are car mechanics.

In this case people are parts of several overlapping networks, networks of Euclidean position, networks of professional relationships, and so forth. Vertices which are distant, because of links in one network, may be much closer in another network. Thus it is by being part of several unrelated networks that enables this search to be successful in relatively few steps. Finding good algorithms under different circumstances is one of the major activities in this area [56]. It has also prompted new attempts to gather experimental information on such networks and searches, for example [57].

4.4. Groups

One area where there has been a great interest in networks, and one which predates the Watts and Stogatz paper, is in areas we might broadly label as anthropology. As Milgram's experiment highlights, data on human interactions is much harder to collect yet it often has immediate relevance to us all. One should not be surprised to find that ideas from graph theory have been borrowed and developed by researchers in this field for many years.

One question that is often asked is, what are the groupings or cohesive blocks in this network? If you ask people or organizations to whom they are connected, to which groups they belong, what alliances they have, the trouble is they may well give you the acceptable answer for their society or business world and not the true one. Indeed they may not even be aware of which links are the most important. Gathering the information in whatever way one can is often challenging in this field, but given the data the idea would be to see if the network structure itself gives us unbiased identifications of the real groupings. With this information many other aspects of the way such system works might become more obvious.

Take for example the work study of families in 14th century Florence by Padgett and Ansell [6]. By expressing the relationships between families as edges in a network, figure 13, a distinctive pattern emerges where Cosimo de Medici was at the centre of the block under his control, whereas his rivals for power had a much more diffuse structure. The suggestion is that it is the very structure of the network around the Medici family which led to Cosimo taking over control of Florence in 1434. So whether created by design or as a reaction to external pressures, if one had studied this network, one might have been able to even predict the success of Cosimo. Note though that these data sets are far smaller than those many physicists discuss. It is completely meaningless to try to label this as a random graph, small world or a scale-free graph, etc., it is just too small. Rather, network based ideas are used to identify small groups and cliques. In this case the ideas of block modelling [59] were used. These ideas have been extended, for example in Moody and White [60]. They study the network of social

²⁶Take the edge which decreases the Manhattan distance to the target by the maximum amount.

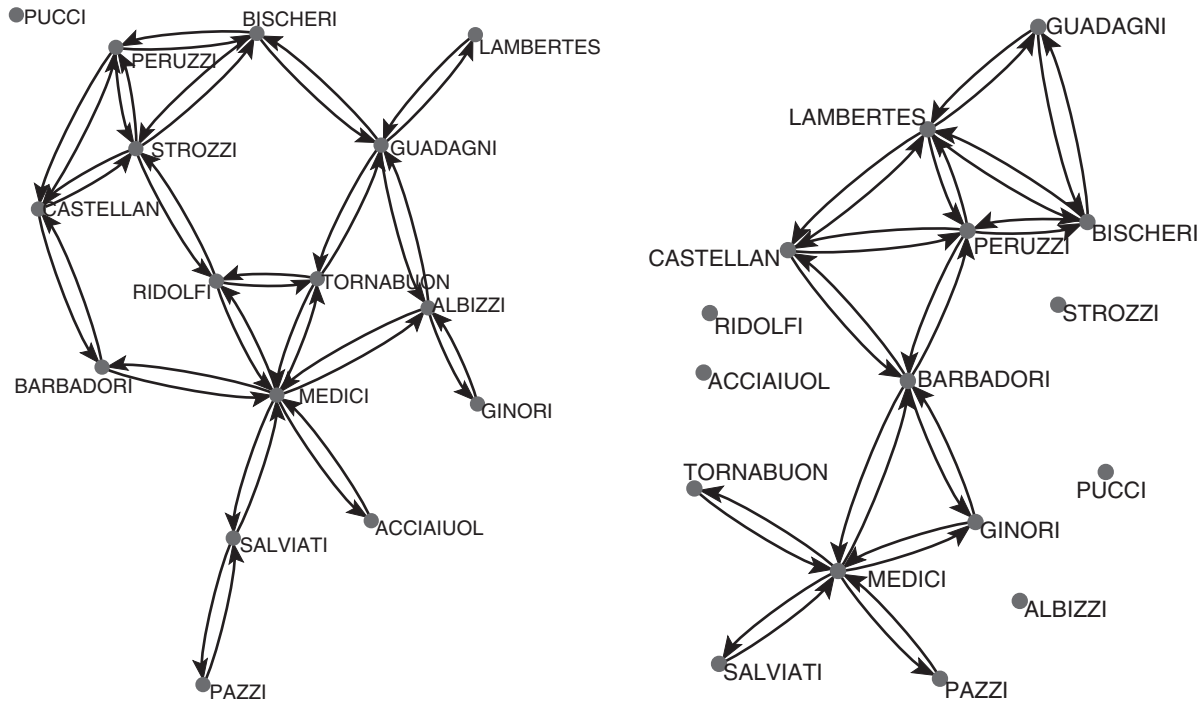


Figure 13. The networks between Florentine families around 1434, marriage on the left, business on the right. Data based on the work of Padgett and taken from examples provided by UCINET [58].

links in an American high school using measures of ‘cohesion’ and ‘embeddedness’ derived from graph-based concepts such as k -connectivity and applying Menger’s theorem. They test their results against ‘outcome variables’. Their groups match the formal organization of the school in terms of year grades but reveal additional details, such as the tendency of younger grades not to be fully assimilated into the central group, and older pupils having more confidence to stay outside the central group, see figure 14.

In a similar way, Hage and Harary [9] used networks to express the different types of information known about the different islands and people in Oceania. These are often based upon exchange of goods, services and information. Again one can gain some understanding of how these cultures operated.

4.5. Dynamical games on a network

Finding a network in the real world, establishing its nature, is often only the precursor to the real question. Often it is not the network but what happens on top of the network that we are really interested in.

For instance, it may be useful to know how computers, switches etc. are connected to form the physical backbone of the internet, but only once we know its structure can we ask how vulnerable is it to random failure or malicious attacks [61]. As we have noted the internet does appear to be scale-free and small world [30]. It was designed by Baran [62] to

have a high level of redundancy so it could withstand a large amount of random damage, because of the original military applications. The many short cuts present in a small world ensure that this is indeed a feature of such networks. If we were to pick a vertex at random and remove it, most paths would lengthen a bit but basically there should be no big problem. Certainly, every vertex would still have a path to all other vertices, the network remains connected. On the other hand, if there was a deliberate attack which removed vertices with the highest degree, the hubs, then a scale-free network will quickly split into several disconnected pieces as so many routes go through a few large hubs [61]. A graph with no hubs, a random or WS small world network, does not suffer unduly under such targeted attacks and still has relatively short paths.

A related topic is virus spreading and immunization, be they computer or human viruses. It is essential for our understanding that we use the right network of contacts when we simulate transmission of viruses. Interestingly, simple models of virus transmission often show a transition from no infection to some infection as immunization levels are lowered (e.g. SIS models). However similar models of virus spreading on scale-free networks often show no such transition. Again, there are too many paths linking one vertex to another [63–65] so that a virus can always find ways of reaching uninfected and unprotected individuals. On the other hand, using the analogy of the malicious attacks on a scale-free network, it is clear that immunising




Figure 14. The network of social links in an American high school analysed by Moody and White using network ideas. Nodes are students and their relationships are links. The vertices are distinguished by their degree k as given in the key. The cohesive groups overlap in ' k -ridges' with components centred on organization by grades. Their interpretation is as follows: 7th graders, core/periphery; 8th, two cliques, one hyper-solidary, the other marginalized; 9th, central transitional; 10th, hang out on margins of seniors; 11th–12th, integrated, but more freedom to marginalize.

the hubs, the vertices with most connections, should bring great dividends. For those whose email is blighted by viruses and spam have probably already realized that immunising random vertices is pretty ineffectual. It only takes a few unprotected computers, a few sites prepared to host the traffic for the infection to spread on a scale-free network. To stop this you have to put the burden on the hubs.

This can be taken much further. Just as with biological virus spreading, there has been interest in all sorts of models of individuals interacting on a small scale to produce large scale effects. This can lead into the realm of simple models such as sand pile models, where large avalanches are produced from simple local rules, or further down the road to agent based modelling in general. In these areas one is looking for power laws, or sudden crashes and changes, an apparent cooperative

macroscopic behaviour in a system with only microscopic rules—a complex system. This for instance has been an area dubbed econophysics (see for example [66, 67]). In these cases the recent advances with networks are merely providing a wider variety of playgrounds to play on. On the other hand if these networks are more realistic, so the network of stock market traders is a small world network, then perhaps the models now have a better chance of fitting the data.

5. Conclusions

I have only been able to scratch the surface of networks. There are many more complicated types of network, many more ways of characterizing them, and far more applications than I have been able to cover here. What is

undoubtedly true is that the last few years have shown that there is a much richer set of possibilities than just the old random graph or lattice, and we have been given many new ways to characterize them. In some ways this is just rationalizing many ideas that had been in existence for some time, often in areas outside physics. Physicists ought to be careful not to assume they are the first to tread along these paths, the roots of networks in other areas goes a long way back. Still, physics can bring its skills and view point along with these new ideas. There are so many applications that there is a lot more mileage in the topic, even if most readily accessible databases have been characterized in terms of networks by now. Modelling real complex systems is inherently difficult, it is often hard to tell if simple models capture the essence of the real system. In many of the problems where networks could be applied, we are working in areas where the data is hard to collect or hard to judge its quality. Trying to study sexually transmitted diseases requires a knowledge of the network of sexual contacts [68]. The network as reported by academic surveys is not going to produce data where we can estimate the errors very easily. Physicists must also be prepared to learn new ways when moving into other areas.

For instance my favourite real world network study is of the social network in Marvel comic characters [69] in which characters are the vertices, linked if they have appeared in the same book. I admit, this was at first because it seemed to be an amusing comment on how far physicists will go to jump on a bandwagon, and perhaps it is no more ludicrous than many other attempts to apply these network ideas. However, other academic areas might find this work of interest even if I do not appreciate it. After all, while we might think that many of our fictional inventions reflect aspects of our life, even those set in some world that has never existed, has any one ever been able to quantify that? Even if as a physicist I was never worried by such a question, it might be that we can help others answer their questions with our new network tools if we can only keep an open mind.

5.1. After this article

The basic ideas and computing tools are not too difficult to pick up. There is now a good mixture of popular sources and more technical books, reviews and papers, and a search for the word network should turn up many more good sources than I have been able to include. The recent book of Watts [18] is a very readable introduction to many topics and has an excellent list of references listed by topics and graded for difficulty. I personally found the earlier book of Watts [22], the review of Dorogovtsev and Mendes [41], the collection of papers edited by Bornholdt and Schuster [70] and Doug White's web site [71] to be excellent sources. There are numerous computer tools and libraries, many free. I have used the JUNG and LEDA libraries [72, 73]

and the tools Pajek and UCINET [58, 74]. These were used to obtain, analyse and display results for many of the models used in the figures here.

Acknowledgements

Jim Moody kindly provided the figure on social links in an American high school, figure 14. The walk algorithm for scale-free networks was developed together with Seb Klauke and Daniel Hook, and the statistical approach to networks generation was discussed with J. B. Laloë and Christian Lunkes. I wish to thank them and several project students for their discussions over the last few years. On a wider scale I have greatly benefited by discussions with members of the ISCOM collaboration and I thank David Lane, Sander van der Leuw and Geoffrey West for providing me with the opportunity to participate as well as many useful conversations. Without ISCOM I would not have been able to have my horizons broadened by Doug White and Carl Knappett and by all the other participants. I also thank the Santa Fe Institute for support during some of these discussions. Finally I would like to thank Daniel Hook, Carl Knappett, Judit Lévai, Ray Rivers and Danièle Steer for comments on the manuscript.

References

- [1] arXiv eprint archive, condensed matter section, <http://arxiv.org/archive/cond-mat>.
- [2] Albert, R., Jeong, H., and Barabási, A. L., 1999, *Nature*, **401**, 130 [eprint: cond-mat/9907038].
- [3] Cancho, R. F. i., and Solé, R. V., 2001, *Proc. R. Soc. (London) B*, **2261**–**2265**, 268.
- [4] Cancho, R. F. i., and Solé, R. V., 2001, *J. Quantit. Linguist.*, **8**, 165 [<http://www.isrl.uiuc.edu/amag/langev/paper/ferrer01twoRegimes.html>].
- [5] Granovetter, M., 1973, *Am. J. Sociol.*, **78**, 1360.
- [6] Padgett, J. F., and Ansell, C. K., 1993, *J. Am. J. Sociol.*, **98**, 1259.
- [7] Broodbank, C., 2000, *An Island Archaeology of the Early Cyclades* (Cambridge: Cambridge University Press).
- [8] Davis, J. L., 1998, *Temple Univ. Aegean Symp.*, **7**, 23.
- [9] Hage, P., and Harary, F., 1991, *Exchange in Oceania: a Graph Theoretic Analysis* (Oxford: Clarendon Press).
- [10] Solomonoff, R., and Rapoport, A., 1950, *Proceedings of the International Congress of Mathematicians*, p. 674.
- [11] Solomonoff, R., and Rapoport, A., 1951, *Bull. Math. Biophys.*, **13**, 107.
- [12] Bollobás, B., 1985, *Random Graphs* (London: Academic Press).
- [13] Erdős, P., and Rényi, A., 1959, *Publ. Math. (Debrecen)*, **6**, 290.
- [14] Bender, C. M., Bender, M. A., Demaine, E. D., and Fekete, S. P., 2004, *J. Phys. A*, **37**, 147.
- [15] Bollobás, B., and Riordan, O. M., 2002, *Handbook of Graphs and Networks—From the Genome to the Internet*, edited by S. Bornholdt and H. G. Stefan (Berlin: Wiley-VCH), p. 1.
- [16] Milgram, S., 1967, *Psychol. Today*, **1**, 60.
- [17] Guare, J., 1999, *Six Degrees of Separation: A Play in Guare: Plays I* (London: Methuen).
- [18] Watts, D. J., 2003, *Six Degrees: The Science of a Connected Age* (London: Heinemann).

- [19] Kleinfeld, J. S., 2002, *Society*, **39**, 61.
- [20] Barabási, A. L., 1999, *Linked: The New Science of Networks* (Cambridge, MA: Perseus Publishing).
- [21] Tjaden, B., 2004, *The Oracle of Bacon*, <http://www.cs.virginia.edu/oracle>.
- [22] Watts, D. J., 1999, *Small Worlds* (Princeton, NJ: Princeton University Press).
- [23] Watts, D. J., and Strogatz, S. H., 1998, *Nature*, **393**, 440.
- [24] Grossman, J., Ion, P., and De Castro, R., 2004, *The Erdős Number Project*, <http://personalwebs.oakland.edu/grossman/erdoshp.html>.
- [25] Lotka, A. J., 1926, *J. Washington Acad. Sci.*, **16**, 317.
- [26] Redner, S., 1998, *Eur. Phys. J. B*, **4**, 131 [e-print: cond-mat/9804163].
- [27] Newman, M. E. J., 2001, *Phys. Rev. E*, **64**, 016131.
- [28] Newman, M. E. J., 2001, *Phys. Rev. E*, **64**, 016132 [e-print: cond-mat/0011144].
- [29] Broder, A., Kumar, R., Maghoul, F., Raghavan, P., Rajagopalan, S., Strata, R., Tomkins, A., and Wiener, J., 2000, *Proceedings of the 9th WWW Conference*, Amsterdam, p. 309. [<http://www9.org/w9cdrom/160/160.html>]
- [30] Faloutsos, M., Faloutsos, P., and Faloutsos, C., 1999, *Comput. Commun. Rev.*, **29**, 251.
- [31] Abello, J., Paradalos, P. M., and Resende, M. G. C., 1999, *DIAMACS Ser. Disc. Math. Theor. Comp. Sci.*, **50**, 119.
- [32] Aiello, W., Chung, F., and Lu, L., 2000, *STOC 00, Proceedings of the Thirty-Second Annual ACM Symposium of Theory of Computing*, pp. 171–180. [<http://doi.acm.org/10.1145/335305.335326>]
- [33] Barabási, A. L., and Albert, R., 1999, *Science*, **286**, 173 [e-print: cond-mat/9910332].
- [34] Jensen, H. J., 1998, *Self-Organized Criticality* (Cambridge: Cambridge University Press).
- [35] Turcotte, D. L., 1999, *Rep. Prog. Phys.*, **62**, 1377.
- [36] US Census Bureau, <http://www.census.gov/census2000/states/us.html>.
- [37] Zipf, G. K., 1949, *Human Behavior and the Principle of Least Effort* (Reading, MA: Addison-Wesley).
- [38] West, G. B., Brown, J. H., and Enquist, B. J., 1997, *Science*, **276**, 122.
- [39] West, G. B., Brown, J. H., and Enquist, B. J., 1999, *Science*, **284**, 1677.
- [40] Asimov, I., 1951, *Foundation* (Trantor: Collins).
- [41] Dorogovtsev, S. N., and Mendes, J. F. F., 2002, *Adv. Phys.*, **51**, 1079 [e-print: cond-mat/0106144].
- [42] Bornholdt, S., and Ebel, H., 2001, *Phys. Rev. E*, **64**, 035104(R) [e-print: cond-mat/008465].
- [43] Simon, H. A., 1955, *Biometrika*, **42**, 425.
- [44] Simon, H. A., 1957, *Models of Man* (New York: Wiley).
- [45] Albert, R., and Barabási, A. L., 2002, *Rev. Mod. Phys.*, **74**, 47 [e-print: cond-mat/0106096].
- [46] Barabási, A. L., Albert, R., and Jeong, H., 1999, *Physica A*, **272**, 173 [e-print: cond-mat/9907068].
- [47] Krapivsky, P. L., Rogers, G. J., and Redner, S., 2001, *Phys. Rev. Lett.*, **86**, 5401 [e-print: cond-mat/0012181].
- [48] Krapivsky, P. L., Redner, S., and Leyvraz, F., 2000, *Phys. Rev. Lett.*, **85**, 4629 [e-print: cond-mat/0005139].
- [49] Klauke, S., 2002, *Scale Free Networks and the Random Walk Algorithm*, University of London Report, April (London: University of London).
- [50] Saramäki, J., and Kaski, K., 2004, e-print: cond-mat/0404088.
- [51] Vázquez, A., 2000, e-print: cond-mat/0006132.
- [52] Vázquez, A., 2003, *Phys. Rev. E*, **67**, 056104 [e-print: cond-mat/0211528].
- [53] Eppstein, D., and Wang, J., 2002, *2nd International Workshop on Web Dynamics*, Honolulu, ACM Computing Research Repository, cs.DM/0204001.
- [54] Adamic, L., Lukose, R. M., and Huberman, B. A., 2002, *Handbook of Graphs and Networks—From the Genome to the Internet*, edited by S. Bornholdt and H. G. Stefan (Berlin: Wiley-VCH), p. 295.
- [55] Kleinberg, J., 2000, *Nature*, **406**, 845.
- [56] Watts, D. J., Dodds, P. S., and Newman, M. E. J., 2002, *Science*, **296**, 1302.
- [57] Dodds, P. S., Muhamed, R., and Watts, D. J., 2003, *Science*, **301**, 827.
- [58] Borgatti, S. P., Everett, M. G., and Freeman, L. C., 2000, *UCINET 6.0 Version 1.00* (Harvard, MA: Analytic Technologies). [<http://www.analytictech.com/ucinet.htm>]
- [59] White, H. C., Boorman, S. A., and Breiger, R. L., 1976, *Am. J. Sociol.*, **81**, 730.
- [60] Moody, J., and White, D. R., 2003, *Am. Sociol. Rev.*, **8**, 1.
- [61] Albert, R., Jeong, H., and Barabási, A. L., 2000, *Nature*, **406**, 379.
- [62] Baran, P., 1964, *Introduction to Distributed Communications Networks*, Report RM = 3420-PR, August [<http://www.rand.org/publications/RM/baran.list.html>].
- [63] Newman, M. E. J., 2002, e-print: cond-mat/0205009.
- [64] Pastor-Satorras, R., and Vespignani, A., 2001, *Phys. Rev. Lett.*, **86**, 3200 [e-print: cond-mat/0010317].
- [65] Pastor-Satorras, R., and Vespignani, A., 2002, *Handbook of Graphs and Networks—From the Genome to the Internet*, edited by S. Bornholdt and H. G. Stefan (Berlin: Wiley-VCH), p. 111 [e-print: cond-mat/0010317].
- [66] Arthur, W. B., 1994, *Am. Econ. Rev.*, **84**, 406.
- [67] Paczuski, M., Bassler, K. E., and Corral, A., 2000, *Phys. Rev. Lett.*, **84**, 3185 [e-print: cond-mat/9905082].
- [68] Liljeros, F., Edling, C., Amaral, L., Stanley, H., and Aberg, Y., 2001, *Nature*, **411**, 907.
- [69] Alberich, R., Miro-Julia, J., and Rosselló, F., 2002, e-print: cond-mat/0202174.
- [70] Bornholdt, S. and Stefan, H. G., (eds), 2002, *Handbook of Graphs and Networks—From the Genome to the Internet* (Berlin: Wiley-VCH).
- [71] Douglas White's web site, <http://eclectic.ss.uci.edu/~drwhite>.
- [72] White, S., Fisher, D., and O'Madadhain, J. *The Java Universal Network/Graph Framework*, <http://jung.sourceforge.net/>.
- [73] Algorithmic Solutions Software GmbH, 2004, *LEDA*, <http://www.algorithmic-solutions.com>.
- [74] Batagelj, V., and Mrvar, A., 1998, *Connections*, **21**, 247 [<http://vlado.fmf.uni-lj.si/pub/networks/pajek>].

Tim Evans is a Senior Lecturer in the Theoretical Physics group of Imperial College London. Apart from three years as a researcher in Canada, he has worked in this group since starting his PhD. The majority of his research is on Thermal Field Theory, how to combine the ideas of Quantum Field Theory with those of statistical mechanics and thermodynamics. This is essential when studying the early universe, relativistic heavy ion collisions and condensed matter systems. It is the central role of phase transitions in such work that led him to study transitions in complex systems and hence to the recent work in complex networks.